



Real Demand and Attested Supply

Zach Kelling

Lux Industries

2026

The number. Morpheus has never published a usage number. No session, price, or revenue series exists in any of the 62 repositories [verified]. Today the network routes about zero verified inference sessions per day [verified]. Real revenue is Lido and Aave yield, about \$476K/yr and falling, down from \$3.24M in the first quarter after the 2024 launch [verified]. Emission runs roughly ten to one over it [verified]. The 18 providers on `active.mor.org` largely resell Venice and OpenAI keys [verified].

The problem, in their numbers. On chain a provider is a base URL and a bearer token. The node dispatches every request to a configured `ApiURL` (`ai_engine.go:64`) and POSTs to it (`openai.go:64`). Settlement is `pricePerSecond` times elapsed time (`SessionRouter.sol:264`); the model and the token counts never enter it [verified]. So a datacenter GPU and a resold ChatGPT subscription settle the same, and a 7B served where a 70B was sold pays the same. The capital layer is sound: principal is withdrawable one to one, and 75% of yield funds a real buyback, about \$511K over the trailing year [verified]. The honest problem is narrower. Dilution outruns revenue, and the meter is fakeable.

The fix, named to a repo. Hanzo connects on both sides with no protocol change. As demand, Enso, Hanzo's learned router, adds Morpheus providers to its candidate set and sends real requests to the best-value bid, opening ordinary sessions on Base [verified]. As supply, Hanzo runs a provider node whose backend is one `models-config` entry pointed at `api.hanzo.ai`, its attested OpenAI-compatible endpoint; the node is `Morpheus-Luminer-Node/proxy-router` and nothing in it changes [verified]. Hanzo attests its GPUs on a root of trust it hosts itself: `lux/ai` checks SPDH measurements against local Reference Integrity Manifests, no cloud. The node's own verifier trusts three remote roots instead, NVIDIA's attestation service (`nras_verifier.go:17`), a third-party quote-parse portal (`verifier.go:80`), and a measurement registry on a CDN mirror (`artifacts_registry.go:20`), any of which can 404 or change its answer [verified]. `mord` meters capacity, work, and identity as three quantities and pays work only on counts the consumer also signed; faked work earns zero and a substituted model is slashed. It is built, 40 tests passing [verified].

The outcome, projected. Assumptions stated. Launch reach across `hanzo.chat`, `hanzo.app`, `hanzo/dev`, `lux.exchange`, and `lux/wallet` is 100M [client]. Assume 2% turn monthly active (2M), 2% of those opt into a private-inference tier (40k), and half route (20k) [illustrative]. Supply bounds the near term at 18 providers: 1,000 to 10,000 verified sessions per day [illustrative]. The low case, about 10,000 per day, still exceeds all currently published Morpheus demand, which is about zero [verified]. Route a small fee on verified volume into the buyback Morpheus already runs, and revenue scales with usage, not deposits.

Relevance, and first step. No other counterparty brings both a live 100M-reach consumer funnel and an un-fakeable meter to point at this network. It adopts today as configuration: `postModelBid` is permissionless for an active provider (`Marketplace.sol:63`), so no facet upgrade and no 5-of-9 Safe signature [verified]. Demand and verifiability land first; governance, listing, and lending rails second; an L2 and a raise third, cleared only into demonstrated on-chain revenue and a KYC base already onboarded through Hanzo IAM and `lux/wallet`.

Diamond 0x6aBE1d282f72B474E54527D93b979A4f64d3030a on Base. Every citation checked against the working tree.

• `research@lux.network`